

# Socket Programming

Kameswari Chebrolu

Reference: Beej's Guide to Network Programming

# Quote

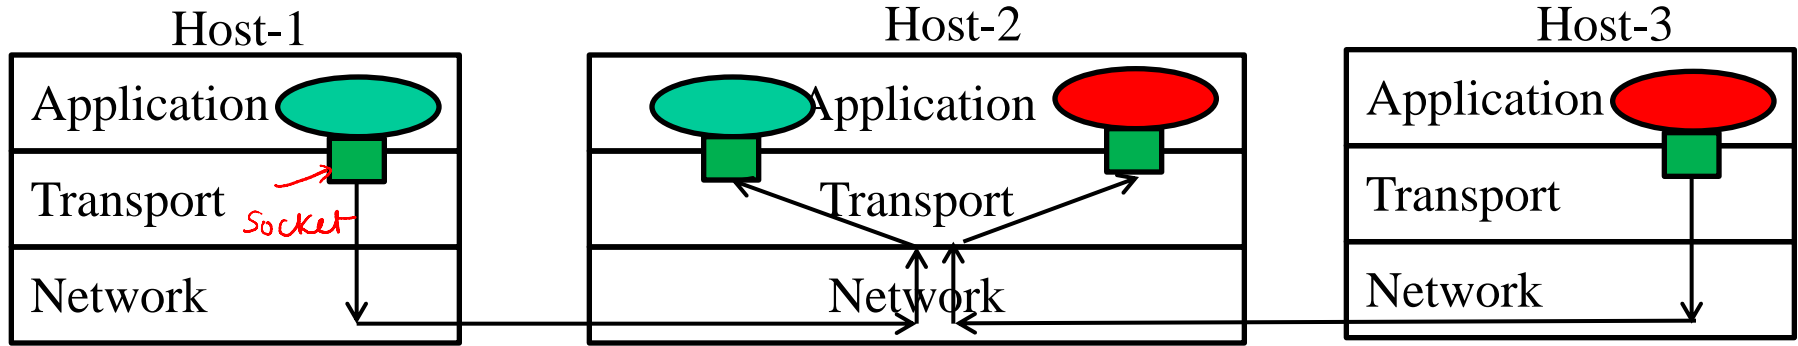
I hear and I forget

I see and I remember

I do and I understand

-- Chinese Proverb

# Multiplexing/Demultiplexing



Demultiplexing: Deliver segments to the right socket

Multiplexing: Assemble segments such that they get delivered to right socket

src & Dest  
IP  
==

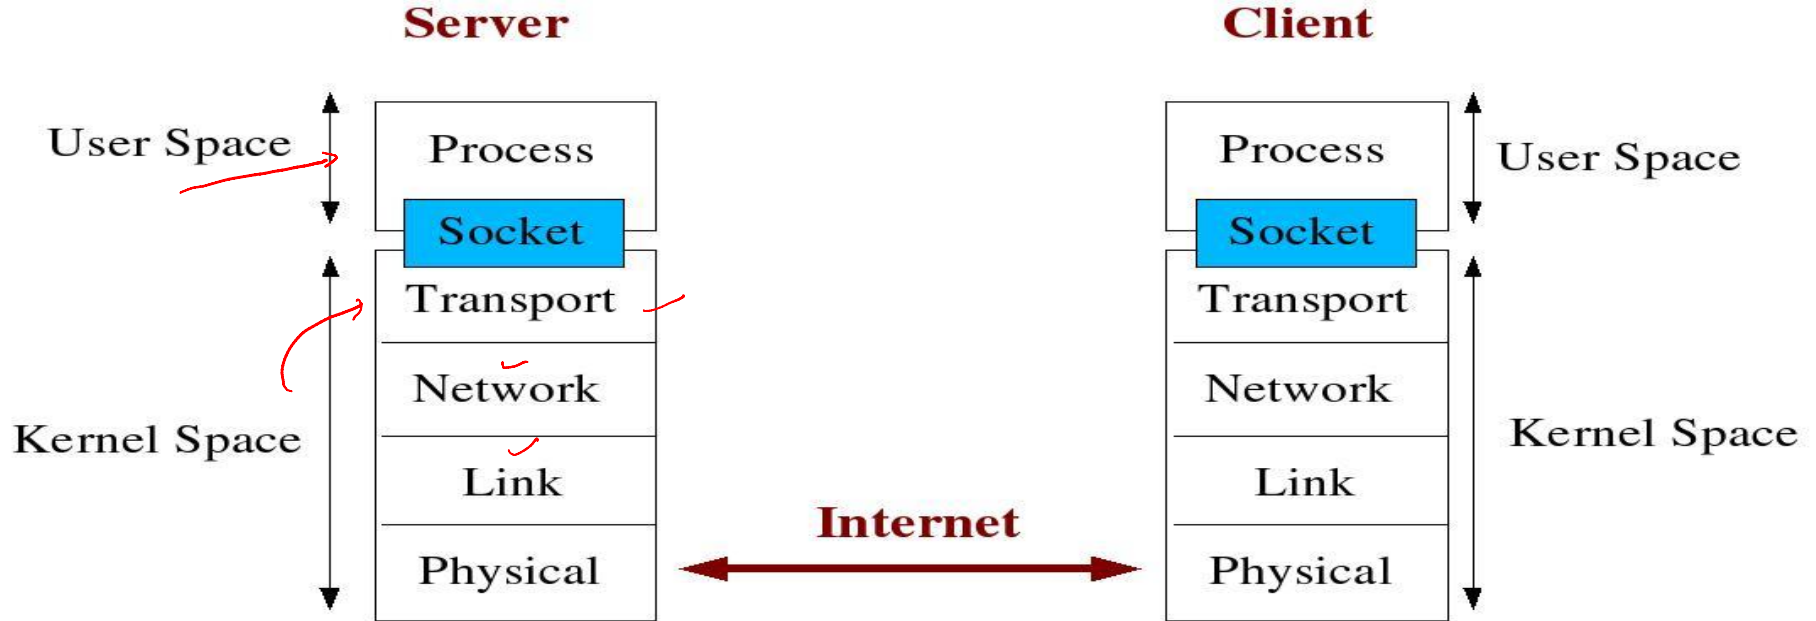
Source Port	Destination Port
Other fields in header	
Application Data	

Transport Layer Segment

# What is a socket?

- Socket: An interface between an application process and transport layer
  - The application process can send/receive messages to/from another application process (local or remote) via a socket
- In Unix jargon, a socket is a file descriptor – an integer associated with an open file
- Types of Sockets: **Internet Sockets**, unix sockets, X.25 sockets etc
  - Internet sockets characterized by IP Address (4 bytes), port number (2 bytes)

# Socket Description



# Types of Internet Sockets

- Stream Sockets (SOCK\_STREAM)
  - Connection oriented
  - Rely on TCP to provide reliable two-way connected communication
- Datagram Sockets (SOCK\_DGRAM)
  - Rely on UDP
  - Connection is unreliable

# Byte Ordering

- Two types of “Byte ordering”
  - Big-Endian (Network Byte Order): High-order byte of the number is stored in memory at the lowest address
  - Little-Endian: Low-order byte of the number is stored in memory at the lowest address
    - Some hosts use this ordering
  - Network stack (TCP/IP) expects Network Byte Order

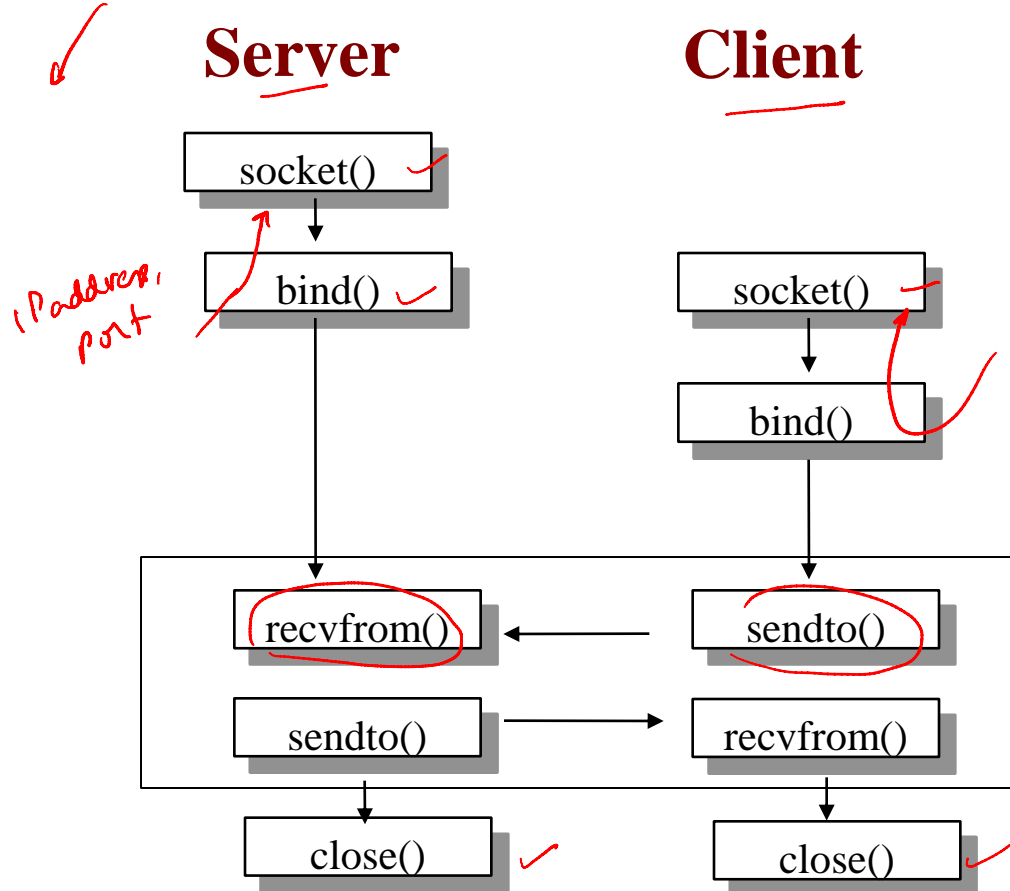
# Byte Ordering

- Conversions:
  - htons() - Host to Network Short
  - htonl() - Host to Network Long
  - ntohs() - Network to Host Short
  - ntohl() - Network to Host Long



# Connectionless Protocol

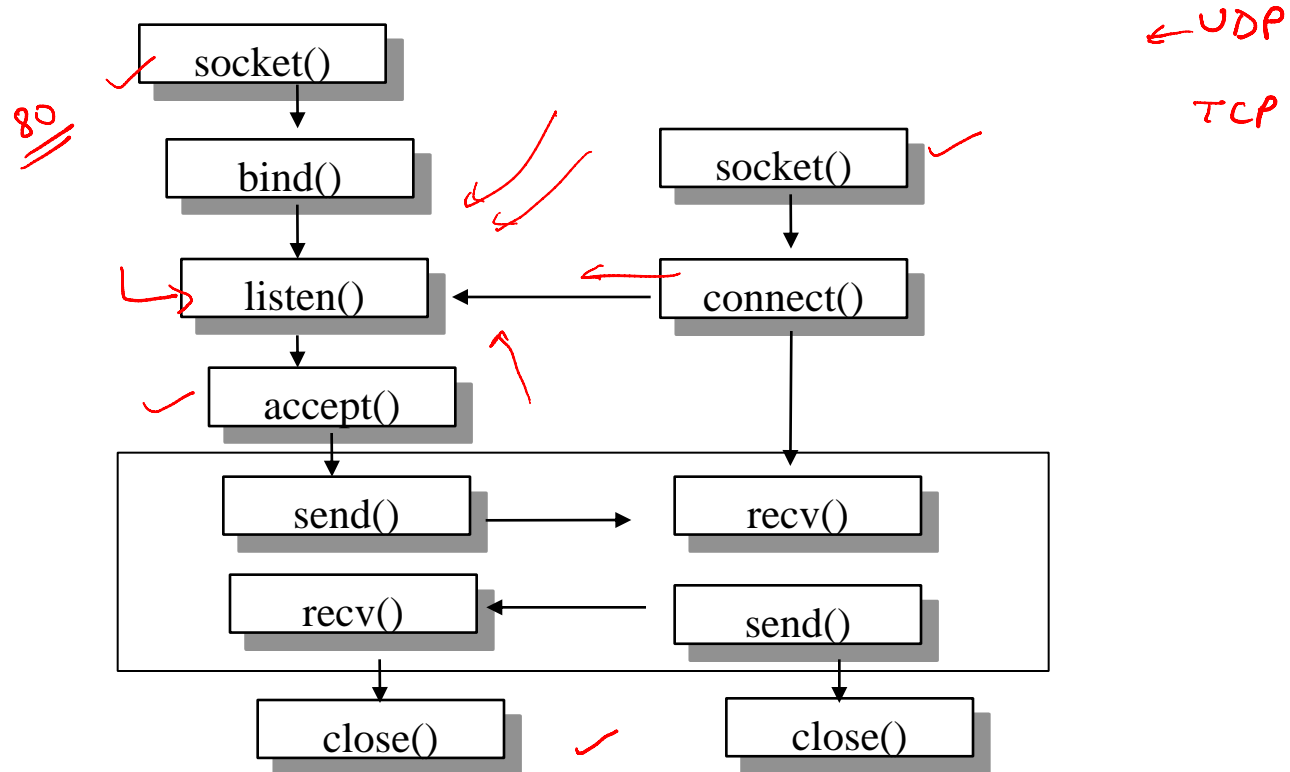
UDP



# Connection Oriented Protocol

Server

Client



# socket() -- Get the file descriptor

- int socket(int domain, int type, int protocol);
  - domain should be set to PF\_INET
  - type can be SOCK\_STREAM or SOCK\_DGRAM  
↳ TCP↳ UDP
  - set protocol to 0 to have socket choose the correct protocol based on type
  - socket() returns a socket descriptor for use in later system calls or -1 on error

```
int sockfd;
```

```
sockfd = socket (PF_INET, SOCK_STREAM, 0);
```

# bind() - what port am I on?

- Used to associate a socket with a port on the local machine
  - The port number is used by the kernel to match an incoming packet to a process)
- int bind(int sockfd, struct sockaddr \*my\_addr, int addrlen)
  - sockfd is the socket descriptor returned by socket()
  - my\_addr is pointer to struct sockaddr that contains information about your IP address and port
  - addrlen is set to sizeof(struct sockaddr)
  - returns -1 on error

# bind() - failure

- All ports below 1024 are reserved
- You can use ports above 1024 upto 65535 provided there are not already in use
- Re-running a server may result in bind failure
  - Why? Socket still around in kernel using the port
  - Solution: Wait a minute or two or use function setsockopt() to clear the socket

# Socket Structures

- struct sockaddr: Holds socket address information for many types of sockets

```
struct sockaddr {  
    ✓ unsigned short  sa_family;    //address family AF_XXXINET  
    ✓ unsigned short  sa_data[14]; //14 bytes of protocol addr  
}
```

# Socket Structures

- struct sockaddr\_in: A parallel structure that makes it easy to reference elements of the socket address

```
struct sockaddr_in {  
    short int  
    ✓ unsigned short int  
    struct in_addr  
    unsigned char  
}
```

*host byte order*

```
    sin_family;    // set to AF_INET  
    sin_port; - 2  // Port number  
    sin_addr; → 4 // Internet address  
    sin_zero[8]; //set to all zeros
```

- sin\_port and sin\_addr must be in **network byte order**

# Populating the structure

```
struct in_addr {  
    unsigned long s_addr; // that's 32-bit long, or 4 bytes  
};
```

- `int inet_aton(const char *cp, struct in_addr *inp);`

*#define MYPORT 80*

```
struct sockaddr_in my_addr;  
my_addr.sin_family = AF_INET; ✓  
my_addr.sin_port = htons(MYPORT);  
inet_aton("10.0.0.5", &(my_addr.sin_addr));  
memset(&(my_addr.sin_zero), '\0', 8);
```

– `inet_aton()` gives non-zero on success; zero on failure



- To convert binary IP to string: `inet_ntoa()`  
`printf("%s", inet_ntoa(my_addr.sin_addr));`
- `my_addr.sin_port = 0;` //choose an unused port at random
- `my_addr.sin_addr.s_addr = INADDR_ANY;` //use my IP adr

# Example

int sockfd;

struct sockaddr\_in my\_addr;

sockfd = socket(PF\_INET, SOCK\_STREAM, 0);

my\_addr.sin\_family = AF\_INET; // host byte order

my\_addr.sin\_port = htons(MYPORT); // short, network byte order

my\_addr.sin\_addr.s\_addr = inet\_addr("10.0.0.1");

memset(&(my\_addr.sin\_zero), '\0', 8); // zero the rest of the struct

bind(sockfd, (struct sockaddr \*)&my\_addr, sizeof(struct sockaddr));

/\*\*\*\*\*\* Code needs error checking. Don't forget to do that \*\*\*\*\*/

# sendto() and recvfrom() - DGRAM style


- UDP SOCK-DGRAM
- int sendto(int sockfd, const void \*msg, int len, int flags, const struct sockaddr \*to, int tolen);

- ↳ 600 bytes
- sockfd: socket descriptor you want to send data to
  - msg is pointer to the data you want to send
  - to is a pointer to a struct sockaddr which contains the destination IP and port
  - tolen is sizeof(struct sockaddr)
  - Set flags to zero
  - Function returns the number of bytes actually sent or -1 on error
- ↳ 500 bytes → 100  
↳ 600 bytes

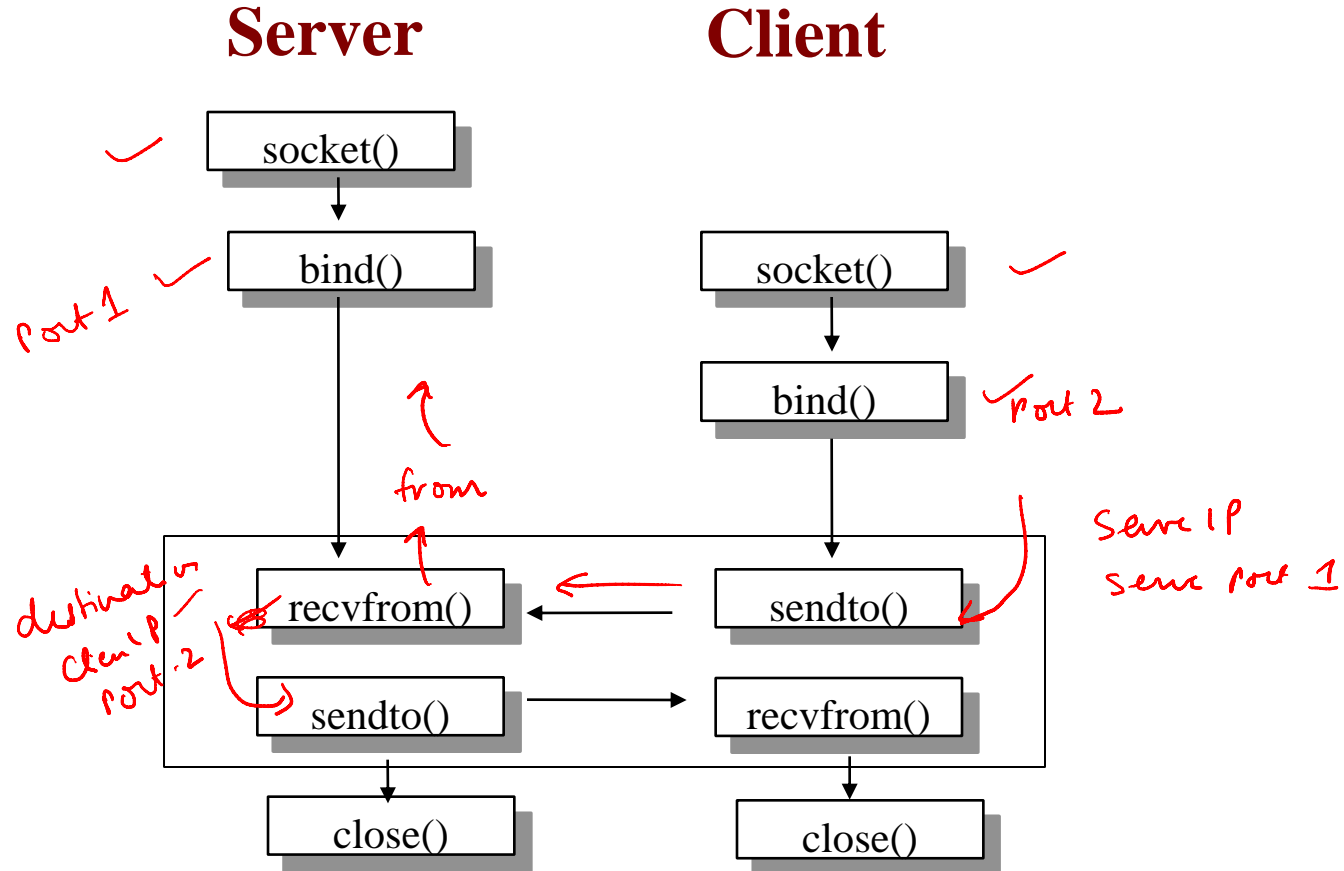
# sendto() and recvfrom() - DGRAM style

- `int recvfrom(int sockfd, void *buf, int len, int flags, struct sockaddr *from, int *fromlen);`
  - *sockfd*: socket descriptor to read from
  - *buf*: buffer to read the information from
  - *len*: maximum length of the buffer
  - *flags* set to zero
  - *from* is a pointer to a local struct sockaddr that will be filled with IP address and port of the originating machine
  - *fromlen* will contain length of address stored in *from*
  - Returns the number of bytes received or -1 on error

# close() - Bye Bye!

- int close(int sockfd);
  - Closes connection corresponding to the socket descriptor and frees the socket descriptor
  - Will prevent any more sends and receives

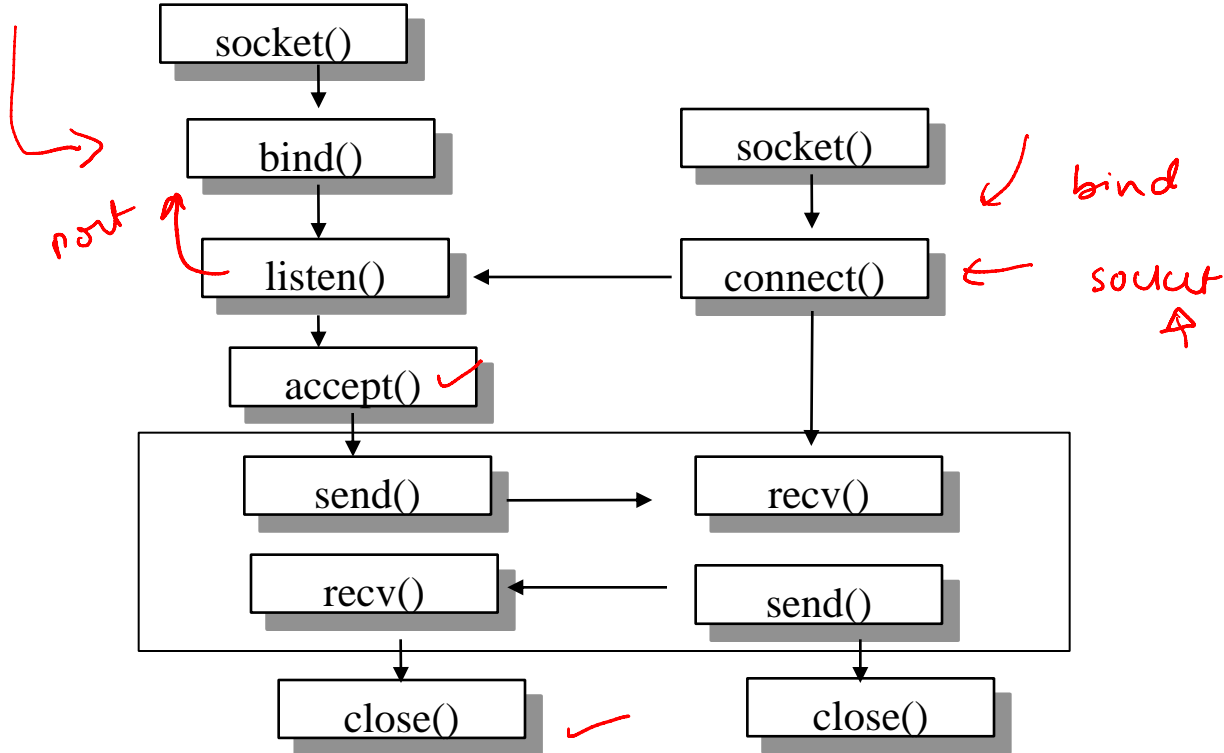
# Connectionless Protocol



# Connection Oriented Protocol

# Server

# Client



# Break

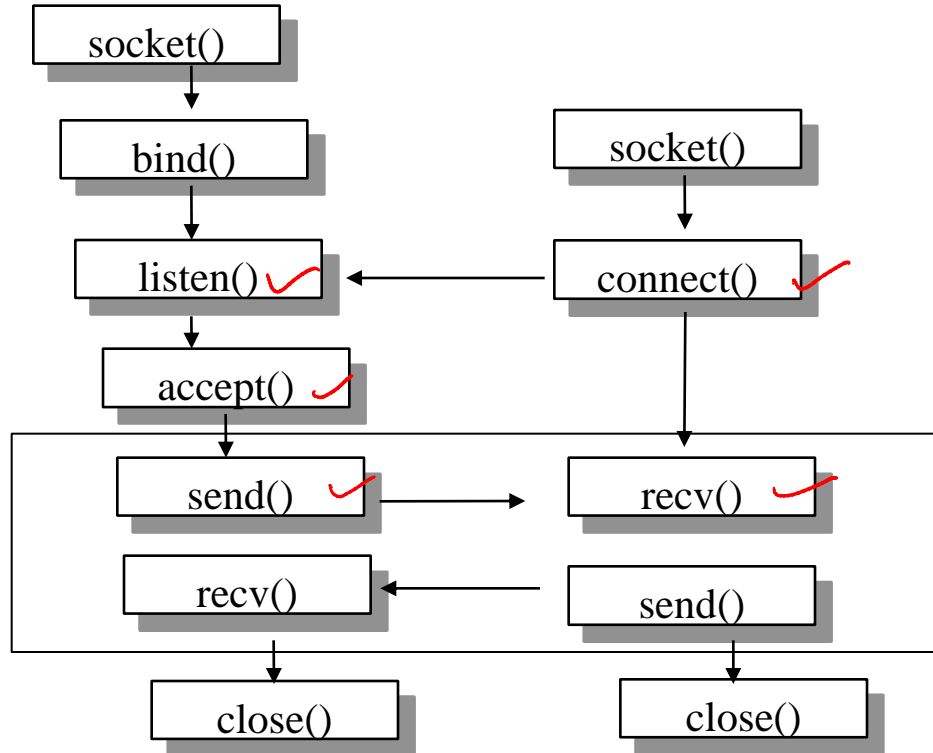




# Connection Oriented Protocol

**Server**

**Client**



# connect() - Hello!

- Connects to a remote host
- int connect(int sockfd, struct sockaddr \*serv\_addr, int addrlen)
  - sockfd is the socket descriptor returned by socket()
  - serv\_addr is pointer to struct sockaddr that contains information on destination IP address and port
  - addrlen is set to sizeof(struct sockaddr)
  - returns -1 on error
- No need to bind(), kernel will choose a port

src, src IP  
random

tuple  
↳ socket

# Example

```
#define DEST_IP  "10.2.44.57"
```

```
#define DEST_PORT 5000
```

```
main(){
```

```
    int sockfd;
```

```
    struct sockaddr_in dest_addr;  // will hold the destination addr
```

```
    sockfd = socket(PF_INET, SOCK_STREAM, 0);
```

```
    dest_addr.sin_family = AF_INET;      // host byte order
```

```
    dest_addr.sin_port = htons(DEST_PORT);  // network byte order
```

```
    dest_addr.sin_addr.s_addr = inet_addr(DEST_IP);
```

```
    memset(&(dest_addr.sin_zero), '\0', 8); // zero the rest of the struct
```

```
    connect(sockfd, (struct sockaddr *)&dest_addr, sizeof(struct sockaddr));
```

```
    /***** Don't forget error checking *****/
```

✓ client random src port

# listen() - Call me please!

- Waits for incoming connections
- `int listen(int sockfd, int backlog);`
  - sockfd is the socket file descriptor returned by `socket()`
  - backlog is the number of connections allowed on the incoming queue
  - `listen()` returns -1 on error
  - Need to call bind() before you can listen()

# accept() - Thank you for calling !

- accept() gets the pending connection on the port you are listen()ing on
- int accept(int sockfd, void \*addr, int \*addrlen);
  - sockfd is the listening socket descriptor
  - information about incoming connection is stored in addr which is a pointer to a local struct sockaddr\_in
  - addrlen is set to sizeof(struct sockaddr\_in)
  - accept returns a new socket file descriptor to use for this accepted connection and -1 on error *↪ send/recv data*

# Example

```
#include <string.h>
```

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

*server side*

```
#include <netinet/in.h>
```

```
#define MYPORT 3490 // the port users will be connecting to
```

```
#define BACKLOG 10 // pending connections queue will hold
```

```
main(){
```

```
    int sockfd, new_fd; // listen on sockfd, new connection on new_fd
```

```
    struct sockaddr_in my_addr; // my address information
```

```
    struct sockaddr_in their_addr; // connector's address information
```

```
    int sin_size;
```

✓ TCP

```
    sockfd = socket(PF_INET, SOCK_STREAM, 0);
```


```
my_addr.sin_family = AF_INET;      // host byte order
my_addr.sin_port = htons(MYPORT); // short, network byte order
my_addr.sin_addr.s_addr = INADDR_ANY; // auto-fill with my IP
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct

// don't forget your error checking for these calls:

bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct sockaddr));
listen(sockfd, BACKLOG);

sin_size = sizeof(struct sockaddr_in);

new_fd = accept(sockfd, (struct sockaddr *)&their_addr, &sin_size);
```



client info

# send() and recv() - Let's talk!

- The two functions are for communicating over stream sockets or connected datagram sockets.  
*bind, listen, accept*
- `int send(int sockfd, const void *msg, int len, int flags);`
  - sockfd is the socket descriptor you want to send data to (got from `accept()`)
  - msg is a pointer to the data you want to send
  - len is the length of that data in bytes
  - set flags to 0 for now
  - `send()` returns the number of bytes actually sent (may be less than the number you told it to send) or -1 on error



# Example

```
char *msg = “hello!”;
```

```
int len, bytes_sent;
```

```
..... ✓ ]
```

```
len = strlen(msg);
```

```
bytes_sent = send(sockfd, msg, len, 0);
```



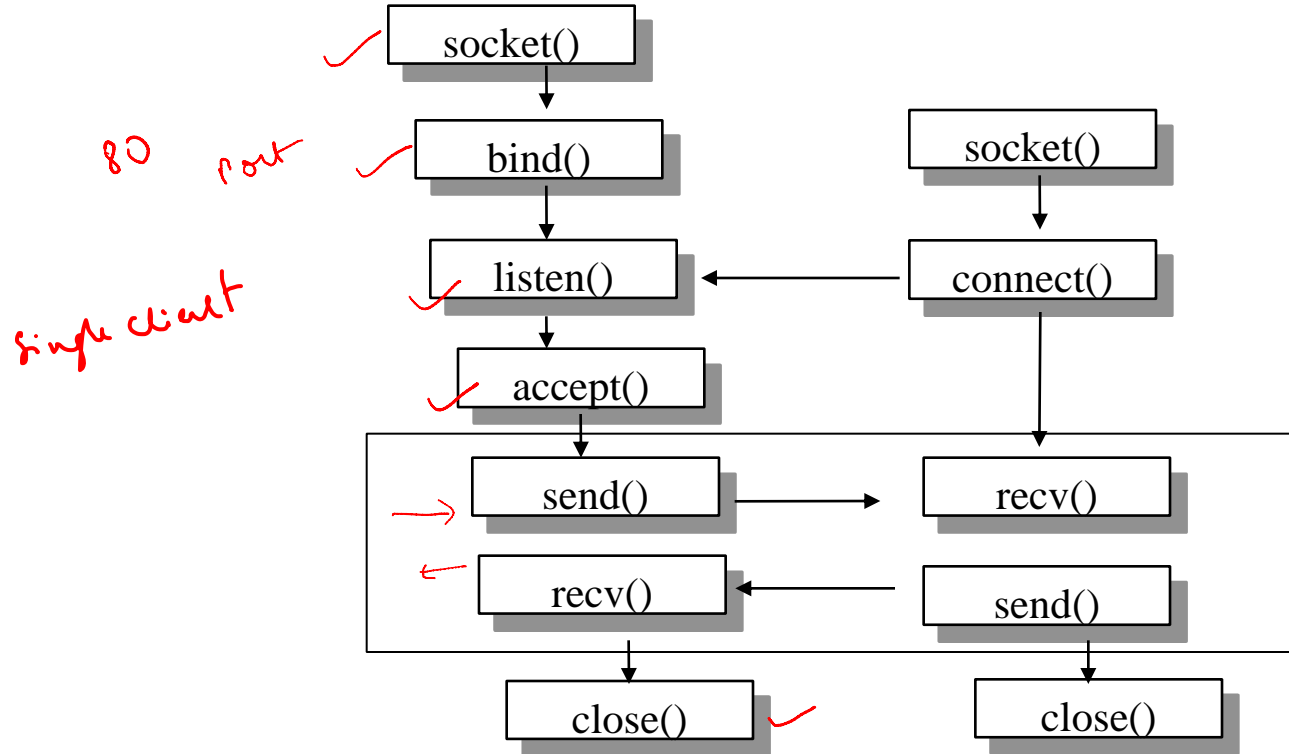
# send() and recv() - Let's talk!

- `int recv(int sockfd, void *buf, int len, int flags);`
  - sockfd is the socket descriptor to read from
  - buf is the buffer to read the information into
  - len is the maximum length of the buffer
  - set flags to 0 for now
  - `recv()` returns the number of bytes actually read into the buffer or -1 on error
  - If `recv()` returns 0, the remote side has closed connection on you

# Connection Oriented Protocol

Server


Client



# Break



# Miscellaneous Routines

- int getpeername(int sockfd, struct sockaddr \*addr, int \*addrlen);  


- Will tell who is at the other end of a connected stream socket and store that info in *addr*

- int gethostname(char \*hostname, size\_t size);  


- Will get the name of the computer your program is running on and store that info in hostname

# Miscellaneous Routines

- Provides DNS service: `struct hostent *gethostbyname(const char *name);`

```
struct hostent {  
    char *h_name;           //official name of host  
    char **h_aliases;       //alternate names for the host  
    int h_addrtype;         //usually AF_INET  
    int h_length;           //length of the address in bytes  
    char **h_addr_list;     //array of network addresses for the host  
}  
#define h_addr h_addr_list[0]
```

- Example Usage:

```
struct hostent *h;  
h = gethostbyname("www.iitb.ac.in");  
printf("Host name : %s \n", h->h_name);  
printf("IP Address: %s\n", inet_ntoa(*((struct in_addr *)h->h_addr)));
```

# Input/Output Multiplexing

client 1 ✓  
client 2 ✓

- Some routines like accept(), recv() block

- Make sockets non-blocking

```
sockfd = socket(PF_INET, SOCK_STREAM, 0);  
fcntl(sockfd, F_SETFL, O_NONBLOCK);
```

- Polling (consumes CPU time)

- Fork a separate process for each I/O channel ✓
- Threading ✓
- Select system call (HIGHLY RECOMMENDED)

↳ listen() ✓  
✓ accept() ✓  
✓ recv() ↑

# Select()

→ list of file descriptors

list → listen, client 1, client 2  
sockets

- `int select(int numfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);`

- `numfds`: highest file descriptor + 1
- `Readfds`, `writefds`, `exceptfds`: set of file descriptors to monitor for read, write and exception operations
- When `select()` returns, the set of file descriptors is modified to reflect the one that is currently ready
- Timeout: select returns after this period if it still hasn't found any ready file descriptors

```
struct timeval {  
    int tv_sec; // seconds  
    int tv_usec; // microseconds  
};
```



# Useful Macros

- FD\_ZERO(fd\_set \*set)
  - clears a file descriptor set
- FD\_SET(int fd, fd\_set \*set)
  - adds fd to the set
- FD\_CLR(int fd, fd\_set \*set)
  - removes fd from the set
- FD\_ISSET(int fd, fd\_set \*set)
  - tests to see if fd is in the set

# Example

```
#define STDIN 0 // file descriptor for standard input
```

```
int main(void) {
```

```
    struct timeval tv;
```

```
    fd_set readfds;
```

```
    tv.tv_sec = 2;
```

```
    tv.tv_usec = 500000;
```

```
    FD_ZERO(&readfds);
```

```
    FD_SET(STDIN, &readfds);
```

```
    // don't care about writefds and exceptfds:
```

```
    select(STDIN+1, &readfds, NULL, NULL, &tv);
```

## Example Cont....

```
if (FD_ISSET(STDIN, &readfds))  
    printf("A key was pressed!\n");  
else  
    printf("Timed out.\n");  
return 0;  
}
```

# Summary

- Sockets help application process to communicate with each other using standard Unix file descriptors
- Two types of Internet sockets: SOCK\_STREAM <sup>→ TCP</sup> and SOCK\_DGRAM <sup>→ UDP</sup>
- Many routines exist to help ease the process of communication

# References

- Books:

- Unix Network Programming, volumes 1-2 by W. Richard Stevens.
- TCP/IP Illustrated, volumes 1-3 by W. Richard Stevens and Gary R. Wright

- Web Resources:

- Beej's Guide to Network Programming

(These slides followed 2001 version, there is a 2012 version that includes IPv6)